

Logistic regression (OvR) - BoW

May 7, 2024

1 Initialization

This notebook will train the Logistic Regression in **One vs Rest** decision function. The Multinomial Logistic Regression is in the Softmax Regression notebook

Connect to Google Drive:

```
[ ]: # from google.colab import drive
      # drive.mount('/content/drive')

      # %cd '/content/drive/MyDrive/GitHub/emotion-detection-from-text'
```

Preparing necessary packages (may need to add more):

```
[ ]: import numpy as np
      import matplotlib.pyplot as plt
      import seaborn as sns
      import pandas as pd
      import warnings
      warnings.filterwarnings('ignore')

      from sklearn.linear_model import LogisticRegression
      from sklearn.model_selection import GridSearchCV, cross_val_score
      from sklearn.metrics import accuracy_score
      from sklearn.preprocessing import StandardScaler
      from joblib import dump, load

      from preset_function import evaluate_model, draw_learning_curve, \
          ↪load_processed_data

      X_train_bow, X_test_bow, X_train_tfidf, X_test_tfidf, \
          X_train_bow_L1, X_test_bow_L1, X_train_tfidf_L1, X_test_tfidf_L1 = \
          ↪load_processed_data('input')

      y_train, y_test = load_processed_data('output')

      %matplotlib inline
```

Select dataset:

```
[ ]: X_train = X_train_bow
     X_test = X_test_bow
```

2 Basic training

```
[ ]: lr_model = LogisticRegression(multi_class='ovr')
     lr_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(multi_class='ovr')
```


```
[ ]: evaluate_model(lr_model, X_train, X_test, y_train, y_test,
    ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9848
- Micro F1 score: 0.9848
- Macro F1 score: 0.9816


Score of on test are:

- Accuracy score: 0.8760
- Micro F1 score: 0.8760
- Macro F1 score: 0.8411





```
[ ]: draw_learning_curve(lr_model, X_train, y_train)
```



3 Multiple tuning

3.1 L1 regularization

First, we try to plot the validation score through a list of C from 0.001 to 100

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l1', solver='liblinear',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))
```


```
trs_list.append(train_score)
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.001, 0.01, 0.1, 1, 5, 10, 100]
[0.3390625, 0.3388125, 0.864625, 0.926625, 0.991375, 0.9924375, 0.9926875]
[0.3368125, 0.3386875, 0.8408749999999999, 0.8725624999999999,
0.8686875000000001, 0.8680625, 0.8635624999999999]

[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



We can see the good value of C is near $C = 1$, then we scope to $C = 1$:

```
[ ]: C_list = [0.1, 0.25, 0.5, 1, 1.25, 1.5, 1.75, 2]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l1', solver='liblinear',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))


    trs_list.append(train_score)
    cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.1, 0.25, 0.5, 1, 1.25, 1.5, 1.75, 2]
[0.864625, 0.88325, 0.8998125, 0.927, 0.939125, 0.953, 0.9664375, 0.984875]
[0.8407500000000001, 0.8659375, 0.8708750000000001, 0.8723749999999999,
0.8714375000000001, 0.8708125000000001, 0.8696249999999999, 0.8695625]

[ ]: [Text(0, 0, '0.1'),
      Text(1, 0, '0.25'),
      Text(2, 0, '0.5'),
      Text(3, 0, '1'),
      Text(4, 0, '1.25'),
      Text(5, 0, '1.5'),
      Text(6, 0, '1.75'),
      Text(7, 0, '2')]
```



We choose $C = 1$ to be the best one

```
[ ]: best_l1_lr_model = LogisticRegression(C=1, penalty='l1', solver='liblinear',  
    ↪multi_class='ovr')
```


```
[ ]: best_l1_lr_model.fit(X_train, y_train)  
    evaluate_model(best_l1_lr_model, X_train, X_test, y_train, y_test,  
    ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9267
- Micro F1 score: 0.9267
- Macro F1 score: 0.9077


Score of on test are:

- Accuracy score: 0.8782
- Micro F1 score: 0.8782
- Macro F1 score: 0.8457





```
[ ]: draw_learning_curve(best_l1_lr_model, X_train, y_train)
```



3.2 L2 regularization

We do the same things from here

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
```

```
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```


```
[0.001, 0.01, 0.1, 1, 5, 10, 100]
```

```
[0.3545, 0.6278125, 0.928, 0.9848125, 0.9916875, 0.9924375, 0.9926875]
```

```
[0.3431875, 0.5831875, 0.8396250000000001, 0.8709374999999999,
```

```
0.8725625000000001, 0.8714375000000001, 0.8636250000000001]
```

```
[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



It looks like good C is still near 1

```
[ ]: C_list = [0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪n_jobs=8))

    trs_list.append(train_score)
    cvs_list.append(cv_score)


[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2]
[0.9606875, 0.97575, 0.981875, 0.9848125, 0.9866875, 0.98825, 0.988875,
0.9894375]
[0.8585, 0.8670000000000002, 0.8691875, 0.8709374999999999, 0.8712499999999999,
0.8717500000000001, 0.8720625, 0.8716250000000001]

[ ]: [Text(0, 0, '0.25'),
      Text(1, 0, '0.5'),
      Text(2, 0, '0.75'),
      Text(3, 0, '1'),
      Text(4, 0, '1.25'),
      Text(5, 0, '1.5'),
```

```
Text(6, 0, '1.75'),  
Text(7, 0, '2')]
```



We choose $C = 1.75$

```
[ ]: best_l2_lr_model = LogisticRegression(C=1.75, penalty='l2', solver='lbfgs',  
↳ multi_class='ovr')
```


```
[ ]: best_l2_lr_model.fit(X_train, y_train)  
evaluate_model(best_l2_lr_model, X_train, X_test, y_train, y_test,  
↳ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9889
- Micro F1 score: 0.9889
- Macro F1 score: 0.9867


Score of on test are:

- Accuracy score: 0.8770
- Micro F1 score: 0.8770
- Macro F1 score: 0.8419





```
[ ]: draw_learning_curve(best_l2_lr_model, X_train, y_train)
```



3.3 Elastic regularization

```
[ ]: dict_param = {
    'C' : [0.001, 0.01, 0.1, 1, 5, 10, 100],
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

lr_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='ovr')
grid_search = GridSearchCV(lr_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)

[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='ovr',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': [0.001, 0.01, 0.1, 1, 5, 10, 100],
                'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```



```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
df = df[df['score'] < 0.8]
print("Bad hyperparameter:")
for param in dict_param:
    for value in dict_param[param]:
        if len(df[df[param] == value]) == 35 // len(dict_param[param]):
            print(param, value)
```

Bad hyperparameter:

C 0.001

C 0.01

```
[ ]: dict_param = {
    'C' : np.logspace(0, 2, 5),
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

lr_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='ovr')
grid_search = GridSearchCV(lr_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='ovr',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': array([ 1.          ,  3.16227766, 10.          ,
31.6227766 ,
100.          ]),
                'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```

```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
print(df)
```

	C	l1_ratio	score
0	1.000000	0.1	0.871063
1	1.000000	0.3	0.872437
2	1.000000	0.5	0.873688
3	1.000000	0.7	0.873250
4	1.000000	0.9	0.873937
5	3.162278	0.1	0.872125
6	3.162278	0.3	0.873375
7	3.162278	0.5	0.874188
8	3.162278	0.7	0.874313
9	3.162278	0.9	0.874375
10	10.000000	0.1	0.871813
11	10.000000	0.3	0.872750
12	10.000000	0.5	0.873625
13	10.000000	0.7	0.873625
14	10.000000	0.9	0.873562
15	31.622777	0.1	0.872500
16	31.622777	0.3	0.872437
17	31.622777	0.5	0.872375
18	31.622777	0.7	0.872188
19	31.622777	0.9	0.872687
20	100.000000	0.1	0.872062
21	100.000000	0.3	0.871937
22	100.000000	0.5	0.872375
23	100.000000	0.7	0.871687
24	100.000000	0.9	0.872125

```
[ ]: print(grid_search.best_estimator_, grid_search.best_score_)
```

```
LogisticRegression(C=3.1622776601683795, l1_ratio=0.9, multi_class='ovr',
                    penalty='elasticnet', solver='saga') 0.874375
```

```
[ ]: best_en_lr_model = LogisticRegression(C=3.1622776601683795, l1_ratio=0.9,
    ↪ multi_class='ovr',
    ↪ penalty='elasticnet', solver='saga')
```


```
[ ]: best_en_lr_model.fit(X_train, y_train)
    evaluate_model(best_en_lr_model, X_train, X_test, y_train, y_test,
    ↪ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9756
- Micro F1 score: 0.9756
- Macro F1 score: 0.9701


Score of on test are:

- Accuracy score: 0.8802
- Micro F1 score: 0.8802
- Macro F1 score: 0.8441





```
[ ]: draw_learning_curve(best_en_lr_model, X_train, y_train)
```



4 Conclusion

There are a few difference among the accuracy of these 3 regularization. However, Elastic-net regularization gives the best performance then I will choose it to be the best model in this notebook.

```
[ ]: best_lr_model = best_en_lr_model
```

```
[ ]: directory = "data/models/lr/"  
  
     dump(best_lr_model, directory + "best_lr_bow_model.joblib")
```

```
[ ]: ['data/models/lr/best_lr_bow_model.joblib']
```