

Logistic regression (OvR) - BoW_L1

May 6, 2024

1 Initialization

This notebook will train the Logistic Regression in **One vs Rest** decision function. The Multinomial Logistic Regression is in the Softmax Regression notebook

Connect to Google Drive:

```
[ ]: # from google.colab import drive
# drive.mount('/content/drive')

# %cd '/content/drive/MyDrive/GitHub/emotion-detection-from-text'
```

Preparing necessary packages (may need to add more):

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import warnings
warnings.filterwarnings('ignore')

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from joblib import dump, load

from preset_function import evaluate_model, draw_learning_curve, \
    load_processed_data

X_train_bow, X_test_bow, X_train_tfidf, X_test_tfidf, \
    X_train_bow_L1, X_test_bow_L1, X_train_tfidf_L1, X_test_tfidf_L1 = \
    load_processed_data('input')

y_train, y_test = load_processed_data('output')

%matplotlib inline
```

Select dataset:

```
[ ]: X_train = X_train_bow_L1
X_test = X_test_bow_L1
```

2 Basic training

```
[ ]: lr_model = LogisticRegression(multi_class='ovr')
lr_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(multi_class='ovr')
```


```
[ ]: evaluate_model(lr_model, X_train, X_test, y_train, y_test,
    ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9682
- Micro F1 score: 0.9682
- Macro F1 score: 0.9576


Score of on test are:

- Accuracy score: 0.8710
- Micro F1 score: 0.8710
- Macro F1 score: 0.8334





```
[ ]: draw_learning_curve(lr_model, X_train, y_train)
```



3 Multiple tuning

3.1 L1 regularization

First, we try to plot the validation score through a list of C from 0.001 to 100

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l1', solver='liblinear',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))
```


```
trs_list.append(train_score)
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.001, 0.01, 0.1, 1, 5, 10, 100]
[0.3390625, 0.3388125, 0.8553125, 0.9159375, 0.98325, 0.9850625, 0.9863125]
[0.3368125, 0.3386875, 0.8334999999999999, 0.8641250000000001,
0.8614375000000001, 0.8583125, 0.851]

[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



We can see the good value of C is near $C = 1$, then we scope to $C = 1$:

```
[ ]: C_list = [0.1, 0.25, 0.5, 1, 1.25, 1.5, 1.75, 2]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l1', solver='liblinear',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))


    trs_list.append(train_score)
    cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.1, 0.25, 0.5, 1, 1.25, 1.5, 1.75, 2]
[0.8553125, 0.874375, 0.8895625, 0.9160625, 0.927375, 0.9419375, 0.9563125,
0.9716875]
[0.8334999999999999, 0.857375, 0.86175, 0.8640625, 0.8644375,
0.8633124999999999, 0.8636250000000001, 0.8627500000000001]

[ ]: [Text(0, 0, '0.1'),
      Text(1, 0, '0.25'),
      Text(2, 0, '0.5'),
      Text(3, 0, '1'),
      Text(4, 0, '1.25'),
      Text(5, 0, '1.5'),
      Text(6, 0, '1.75'),
      Text(7, 0, '2')]
```



We choose $C = 1.25$ to be the best one

```
[ ]: best_l1_lr_model = LogisticRegression(C=1.25, penalty='l1', solver='liblinear',  
    ↪ multi_class='ovr')
```


```
[ ]: best_l1_lr_model.fit(X_train, y_train)  
    evaluate_model(best_l1_lr_model, X_train, X_test, y_train, y_test,  
    ↪ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9274
- Micro F1 score: 0.9274
- Macro F1 score: 0.9062


Score of on test are:

- Accuracy score: 0.8675
- Micro F1 score: 0.8675
- Macro F1 score: 0.8274





```
[ ]: draw_learning_curve(best_l1_lr_model, X_train, y_train)
```



3.2 L2 regularization

We do the same things from here

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
```

```
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)


# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[0.001, 0.01, 0.1, 1, 5, 10, 100]
```

```
[0.3533125, 0.622, 0.90475, 0.9681875, 0.9833125, 0.98475, 0.9863125]
```

```
[0.34331249999999996, 0.584625, 0.83762500000000001, 0.86525, 0.86549999999999999,
0.86343750000000001, 0.85406250000000001]
```

```
[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



It looks like good C is still near 1

```
[ ]: C_list = [0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪n_jobs=8))

    trs_list.append(train_score)
    cvs_list.append(cv_score)


[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2]
[0.9371875, 0.9559375, 0.9629375, 0.9681875, 0.9714375, 0.9735, 0.975,
0.9766875]
[0.8557500000000001, 0.8615, 0.8641249999999999, 0.86525, 0.86575, 0.8655625,
0.8650625000000002, 0.864875]

[ ]: [Text(0, 0, '0.25'),
      Text(1, 0, '0.5'),
      Text(2, 0, '0.75'),
      Text(3, 0, '1'),
      Text(4, 0, '1.25'),
      Text(5, 0, '1.5'),
```

```
Text(6, 0, '1.75'),  
Text(7, 0, '2')]
```



We choose $C = 1.25$

```
[ ]: best_l2_lr_model = LogisticRegression(C=1.25, penalty='l2', solver='lbfgs',  
↳ multi_class='ovr')
```


```
[ ]: best_l2_lr_model.fit(X_train, y_train)  
evaluate_model(best_l2_lr_model, X_train, X_test, y_train, y_test,  
↳ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9714
- Micro F1 score: 0.9714
- Macro F1 score: 0.9624


Score of on test are:

- Accuracy score: 0.8712
- Micro F1 score: 0.8713
- Macro F1 score: 0.8354





```
[ ]: draw_learning_curve(best_l2_lr_model, X_train, y_train)
```



3.3 Elastic regularization

```
[ ]: dict_param = {
    'C' : [0.001, 0.01, 0.1, 1, 5, 10, 100],
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

lr_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='ovr')
grid_search = GridSearchCV(lr_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)

[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='ovr',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': [0.001, 0.01, 0.1, 1, 5, 10, 100],
                'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```



```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
print(df)
df = df[df['score'] < 0.85]
print("Bad hyperparameter:")
for param in dict_param:
    for value in dict_param[param]:
        if len(df[df[param] == value]) == 35 // len(dict_param[param]):
            print(param, value)
```

	C	l1_ratio	score
0	0.001	0.1	0.338688
1	0.001	0.3	0.338688
2	0.001	0.5	0.338688
3	0.001	0.7	0.338688
4	0.001	0.9	0.338688
5	0.010	0.1	0.521188
6	0.010	0.3	0.350500
7	0.010	0.5	0.339000
8	0.010	0.7	0.338687
9	0.010	0.9	0.338687
10	0.100	0.1	0.837938
11	0.100	0.3	0.833937
12	0.100	0.5	0.830187
13	0.100	0.7	0.830688
14	0.100	0.9	0.830375
15	1.000	0.1	0.864812
16	1.000	0.3	0.865750
17	1.000	0.5	0.865563
18	1.000	0.7	0.866313
19	1.000	0.9	0.865750
20	5.000	0.1	0.866687
21	5.000	0.3	0.866812
22	5.000	0.5	0.867500
23	5.000	0.7	0.867375
24	5.000	0.9	0.866937
25	10.000	0.1	0.867062
26	10.000	0.3	0.867125
27	10.000	0.5	0.867000
28	10.000	0.7	0.867563
29	10.000	0.9	0.867062
30	100.000	0.1	0.866938
31	100.000	0.3	0.866500

```

32 100.000      0.5 0.866500
33 100.000      0.7 0.866625
34 100.000      0.9 0.866563

```

Bad hyperparameter:

```

C 0.001
C 0.01
C 0.1

```

```

[ ]: dict_param = {
    'C' : np.logspace(0, 2, 5),
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

lr_model = LogisticRegression(penalty='elasticnet', solver='saga',
    multi_class='ovr')
grid_search = GridSearchCV(lr_model, dict_param, scoring='accuracy', cv=5,
    n_jobs=-1)
grid_search.fit(X_train, y_train)

```

```

[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='ovr',
    penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': array([ 1.          ,  3.16227766, 10.          ,
31.6227766 ,
    100.          ]),
    'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')

```

```

[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
print(df)

```

	C	l1_ratio	score
0	1.000000	0.1	0.864875
1	1.000000	0.3	0.865812
2	1.000000	0.5	0.865500
3	1.000000	0.7	0.866375
4	1.000000	0.9	0.865875
5	3.162278	0.1	0.867125
6	3.162278	0.3	0.867125
7	3.162278	0.5	0.867313
8	3.162278	0.7	0.867000

9	3.162278	0.9	0.867250
10	10.000000	0.1	0.867000
11	10.000000	0.3	0.866938
12	10.000000	0.5	0.867188
13	10.000000	0.7	0.867250
14	10.000000	0.9	0.866937
15	31.622777	0.1	0.866938
16	31.622777	0.3	0.866812
17	31.622777	0.5	0.866812
18	31.622777	0.7	0.866938
19	31.622777	0.9	0.867250
20	100.000000	0.1	0.866750
21	100.000000	0.3	0.866750
22	100.000000	0.5	0.866875
23	100.000000	0.7	0.866437
24	100.000000	0.9	0.867062

```
[ ]: print(grid_search.best_estimator_, grid_search.best_score_)
```

```
LogisticRegression(C=3.1622776601683795, l1_ratio=0.5, multi_class='ovr',
                    penalty='elasticnet', solver='saga') 0.8673125
```

```
[ ]: best_en_lr_model = LogisticRegression(C=3.1622776601683795, l1_ratio=0.5,
    ↪multi_class='ovr',
    ↪penalty='elasticnet', solver='saga')
```


```
[ ]: best_en_lr_model.fit(X_train, y_train)
    ↪evaluate_model(best_en_lr_model, X_train, X_test, y_train, y_test,
    ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9646
- Micro F1 score: 0.9646
- Macro F1 score: 0.9531


Score of on test are:

- Accuracy score: 0.8720
- Micro F1 score: 0.8720
- Macro F1 score: 0.8316





```
[ ]: draw_learning_curve(best_en_lr_model, X_train, y_train)
```



4 Conclusion

There are a few difference among the accuracy of these 3 regularization. However, Elastic-net regularization gives the best performance then I will choose it to be the best model in this notebook.

```
[ ]: best_lr_model = best_en_lr_model
```

```
[ ]: directory = "data/models/lr/"  
  
     dump(best_lr_model, directory + "best_lr_bow_l1_model.joblib")
```

```
[ ]: ['data/models/lr/best_lr_bow_l1_model.joblib']
```