

Logistic regression (OvR) - TF_IDF

May 6, 2024

1 Initialization

This notebook will train the Logistic Regression in **One vs Rest** decision function. The Multinomial Logistic Regression is in the Softmax Regression notebook

Connect to Google Drive:

```
[ ]: # from google.colab import drive
# drive.mount('/content/drive')

# %cd '/content/drive/MyDrive/GitHub/emotion-detection-from-text'
```

Preparing necessary packages (may need to add more):

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import warnings
warnings.filterwarnings('ignore')

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from joblib import dump, load

from preset_function import evaluate_model, draw_learning_curve, \
    load_processed_data

X_train_bow, X_test_bow, X_train_tfidf, X_test_tfidf, \
    X_train_bow_L1, X_test_bow_L1, X_train_tfidf_L1, X_test_tfidf_L1 = \
    load_processed_data('input')

y_train, y_test = load_processed_data('output')

%matplotlib inline
```

Select dataset:

```
[ ]: X_train = X_train_tfidf
     X_test = X_test_tfidf
```

2 Basic training

```
[ ]: lr_model = LogisticRegression(multi_class='ovr')
     lr_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(multi_class='ovr')
```


```
[ ]: evaluate_model(lr_model, X_train, X_test, y_train, y_test,
                    include_training=True)
```


Score of on train are:

- Accuracy score: 0.9216
- Micro F1 score: 0.9216
- Macro F1 score: 0.8767


Score of on test are:

- Accuracy score: 0.8183
- Micro F1 score: 0.8183
- Macro F1 score: 0.7390





```
[ ]: draw_learning_curve(lr_model, X_train, y_train)
```



3 Multiple tuning

3.1 L1 regularization

First, we try to plot the validation score through a list of C from 0.001 to 100

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l1', solver='liblinear',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))
```


```
trs_list.append(train_score)
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.001, 0.01, 0.1, 1, 5, 10, 100]
[0.3386875, 0.3386875, 0.520875, 0.8801875, 0.9753125, 0.9925625, 0.992625]
[0.33868750000000001, 0.33868750000000001, 0.38818749999999996,
0.86218749999999999, 0.8720625, 0.87068750000000001, 0.8684999999999998]

[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



We can see the good value of C is near $C = 5$, then we scope to $C = 5$:

```
[ ]: C_list = [4.1, 4.25, 4.5, 5, 5.25, 5.5, 5.75, 6]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l1', solver='liblinear',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))


    trs_list.append(train_score)
    cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[4.1, 4.25, 4.5, 5, 5.25, 5.5, 5.75, 6]
[0.9584375, 0.9616875, 0.96775, 0.9753125, 0.97875, 0.981125, 0.9828125,
0.984375]
[0.8724375, 0.8726875, 0.8727499999999999, 0.8721249999999999,
0.8714999999999999, 0.8715624999999999, 0.8714375000000001, 0.8716250000000001]

[ ]: [Text(0, 0, '4.1'),
      Text(1, 0, '4.25'),
      Text(2, 0, '4.5'),
      Text(3, 0, '5'),
      Text(4, 0, '5.25'),
      Text(5, 0, '5.5'),
      Text(6, 0, '5.75'),
      Text(7, 0, '6')]
```



We choose $C = 4.5$ to be the best one

```
[ ]: best_l1_lr_model = LogisticRegression(C=4.5, penalty='l1', solver='liblinear',  
    ↪ multi_class='ovr')
```


```
[ ]: best_l1_lr_model.fit(X_train, y_train)  
    evaluate_model(best_l1_lr_model, X_train, X_test, y_train, y_test,  
    ↪ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9677
- Micro F1 score: 0.9677
- Macro F1 score: 0.9597


Score of on test are:

- Accuracy score: 0.8762
- Micro F1 score: 0.8762
- Macro F1 score: 0.8420





```
[ ]: draw_learning_curve(best_l1_lr_model, X_train, y_train)
```



3.2 L2 regularization

We do the same things from here

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪ multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
```

```
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```


```
[0.001, 0.01, 0.1, 1, 5, 10, 100]
```

```
[0.3386875, 0.3548125, 0.627375, 0.921625, 0.9884375, 0.9918125, 0.992625]
```

```
[0.33868750000000001, 0.33993749999999995, 0.57475000000000001,
```

```
0.8021874999999999, 0.8504375, 0.8563749999999999, 0.86075]
```

```
[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



It looks like good C is near 100 or beyond

```
[ ]: C_list = [100, 150, 200, 250, 300]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    lr_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪multi_class='ovr')
    lr_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, lr_model.predict(X_train))
    cv_score = np.mean(cross_val_score(lr_model, X_train, y_train, cv=5,
    ↪n_jobs=8))

    trs_list.append(train_score)
    cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)


# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[100, 150, 200, 250, 300]
```

```
[0.992625, 0.992625, 0.992625, 0.992625, 0.992625]
```

```
[0.86075, 0.860125, 0.8616875, 0.8619375, 0.8614374999999999]
```

```
[ ]: [Text(0, 0, '100'),
      Text(1, 0, '150'),
      Text(2, 0, '200'),
      Text(3, 0, '250'),
      Text(4, 0, '300')]
```



We choose $C = 250$

```
[ ]: best_l2_lr_model = LogisticRegression(C=250, penalty='l2', solver='lbfgs',
    ↪ multi_class='ovr')
```


```
[ ]: best_l2_lr_model.fit(X_train, y_train)
    evaluate_model(best_l2_lr_model, X_train, X_test, y_train, y_test,
    ↪ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9926
- Micro F1 score: 0.9926
- Macro F1 score: 0.9906


Score of on test are:

- Accuracy score: 0.8645
- Micro F1 score: 0.8645
- Macro F1 score: 0.8242





```
[ ]: draw_learning_curve(best_l2_lr_model, X_train, y_train)
```



3.3 Elastic regularization

```
[ ]: dict_param = {
    'C' : [0.001, 0.01, 0.1, 1, 5, 10, 100],
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

lr_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='ovr')
grid_search = GridSearchCV(lr_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)

[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='ovr',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': [0.001, 0.01, 0.1, 1, 5, 10, 100],
                'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```



```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
print(df)
df = df[df['score'] < 0.86]
print("Bad hyperparameter:")
for param in dict_param:
    for value in dict_param[param]:
        if len(df[df[param] == value]) == 35 // len(dict_param[param]):
            print(param, value)
```

	C	l1_ratio	score
0	0.001	0.1	0.338688
1	0.001	0.3	0.338688
2	0.001	0.5	0.338688
3	0.001	0.7	0.328937
4	0.001	0.9	0.319312
5	0.010	0.1	0.338688
6	0.010	0.3	0.338688
7	0.010	0.5	0.338688
8	0.010	0.7	0.338688
9	0.010	0.9	0.338688
10	0.100	0.1	0.556750
11	0.100	0.3	0.513687
12	0.100	0.5	0.465438
13	0.100	0.7	0.421687
14	0.100	0.9	0.389500
15	1.000	0.1	0.810875
16	1.000	0.3	0.821937
17	1.000	0.5	0.830187
18	1.000	0.7	0.839438
19	1.000	0.9	0.856437
20	5.000	0.1	0.854688
21	5.000	0.3	0.862250
22	5.000	0.5	0.868125
23	5.000	0.7	0.872062
24	5.000	0.9	0.874500
25	10.000	0.1	0.859125
26	10.000	0.3	0.864750
27	10.000	0.5	0.870062
28	10.000	0.7	0.872687
29	10.000	0.9	0.874375
30	100.000	0.1	0.862625
31	100.000	0.3	0.864062

```

32 100.000      0.5  0.866250
33 100.000      0.7  0.868875
34 100.000      0.9  0.870187

```

Bad hyperparameter:

```

C 0.001
C 0.01
C 0.1
C 1

```

```

[ ]: dict_param = {
    'C' : np.logspace(1, 2, 6),
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

lr_model = LogisticRegression(penalty='elasticnet', solver='saga',
    multi_class='ovr')
grid_search = GridSearchCV(lr_model, dict_param, scoring='accuracy', cv=5,
    n_jobs=-1)
grid_search.fit(X_train, y_train)

```

```

[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='ovr',
    penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': array([ 10.          , 15.84893192, 25.11886432,
39.81071706,
63.09573445, 100.          ]),
    'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')

```

```

[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
print(df)

```

	C	l1_ratio	score
0	10.000000	0.1	0.859125
1	10.000000	0.3	0.864625
2	10.000000	0.5	0.870125
3	10.000000	0.7	0.872750
4	10.000000	0.9	0.874125
5	15.848932	0.1	0.861437
6	15.848932	0.3	0.865500
7	15.848932	0.5	0.870750

8	15.848932	0.7	0.872250
9	15.848932	0.9	0.874188
10	25.118864	0.1	0.862188
11	25.118864	0.3	0.866062
12	25.118864	0.5	0.869250
13	25.118864	0.7	0.871750
14	25.118864	0.9	0.874875
15	39.810717	0.1	0.862187
16	39.810717	0.3	0.866187
17	39.810717	0.5	0.868812
18	39.810717	0.7	0.870937
19	39.810717	0.9	0.872750
20	63.095734	0.1	0.862562
21	63.095734	0.3	0.865375
22	63.095734	0.5	0.868563
23	63.095734	0.7	0.870875
24	63.095734	0.9	0.871312
25	100.000000	0.1	0.862563
26	100.000000	0.3	0.864812
27	100.000000	0.5	0.867062
28	100.000000	0.7	0.868563
29	100.000000	0.9	0.870125

```
[ ]: print(grid_search.best_estimator_, grid_search.best_score_)
```

```
LogisticRegression(C=25.118864315095795, l1_ratio=0.9, multi_class='ovr',
                    penalty='elasticnet', solver='saga') 0.8748750000000001
```

```
[ ]: best_en_lr_model = LogisticRegression(C=25.118864315095795, l1_ratio=0.9,
    ↪multi_class='ovr',
    ↪penalty='elasticnet', solver='saga')
```


```
[ ]: best_en_lr_model.fit(X_train, y_train)
    evaluate_model(best_en_lr_model, X_train, X_test, y_train, y_test,
    ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9926
- Micro F1 score: 0.9926
- Macro F1 score: 0.9906


Score of on test are:

- Accuracy score: 0.8785
- Micro F1 score: 0.8785
- Macro F1 score: 0.8423





```
[ ]: draw_learning_curve(best_en_lr_model, X_train, y_train)
```



4 Conclusion

There are a few difference among the accuracy of these 3 regularization. However, Elastic-net regularization gives the best performance then I will choose it to be the best model in this notebook.

```
[ ]: best_lr_model = best_en_lr_model
```

```
[ ]: directory = "data/models/lr/"  
     dump(best_lr_model, directory + "best_lr_tfidf_model.joblib")
```

```
[ ]: ['data/models/lr/best_lr_tfidf_model.joblib']
```