Random Forest - BoW_L1

May 11, 2024

1 Initialization

Connect to Google Drive:

```
[ ]: # from google.colab import drive
     # drive.mount('/content/drive')

     # %cd '/content/drive/MyDrive/GitHub/emotion-dectection-from-text'
```

Preparing necessary packages:

```
[ ]: import numpy as np
     import pandas as pd
     import seaborn as sns

     from sklearn.ensemble import RandomForestClassifier
     from sklearn.model_selection import GridSearchCV, cross_val_score
     from sklearn.metrics import accuracy_score
     from joblib import dump

     from preset_function import evaluate_model, draw_learning_curve,␣
      ↪load_processed_data

     X_train_bow, X_test_bow, X_train_tfidf, X_test_tfidf, \
         X_train_bow_L1, X_test_bow_L1, X_train_tfidf_L1, X_test_tfidf_L1=␣
      ↪load_processed_data('input')

     y_train, y_test = load_processed_data('output')

     %matplotlib inline
```

1.1 Select dataset

At first, we choose the dataset to be used for training and testing the model.

```
[ ]: X_train = X_train_bow_L1
     X_test = X_test_bow_L1
```

2 Basic training

We define the model with the default parameters and train it.

```
[ ]: RF = RandomForestClassifier()
     RF.fit(X_train , y_train)
```

```
[ ]: RandomForestClassifier()
```

Evaluate this model using a preset function:

```
[ ]: evaluate_model(RF, X_train, X_test, y_train, y_test, include_training=True)
```

```
Score of on train are:
        - Accuracy score: 0.9864
        - Micro F1 score: 0.9864
        - Macro F1 score: 0.9825
Score of on test are:
        - Accuracy score: 0.8658
        - Micro F1 score: 0.8658
        - Macro F1 score: 0.8291
```

Confusion matrix on Test

Draw learning curve using a preset function:

```
[ ]: draw_learning_curve(RF, X_train, y_train)
```

3 Single tuning

This section examines the best range for each parameters by plotting the performance of the model with a range of value for each parameters.

3.1 N_estimator

The number of trees in the forest.

```python
# Setting the possible value for n_estimators
n_estimators_list = [32, 64, 128, 256, 512]

trs_list = list()
cvs_list = list()

for n_estimators in n_estimators_list:
    # Define model for each n_estimators
    rf_model = RandomForestClassifier(n_estimators=n_estimators)
    rf_model.fit(X_train, y_train)

    # Calculate the cross validation score
    train_score = accuracy_score(y_train, rf_model.predict(X_train))
```

```
    cvs_score = np.mean(cross_val_score(rf_model, X_train, y_train, cv=5,␣
 ↪n_jobs=-1))

    trs_list.append(train_score)
    cvs_list.append(cvs_score)
```

```python
# Draw the plot for n_estimators
fig = sns.lineplot(x=list(range(len(n_estimators_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(n_estimators_list))), y=trs_list)
fig.set_xticks(range(len(n_estimators_list)))
fig.set_xticklabels(n_estimators_list)
```

```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```

```
[ ]: [Text(0, 0, '32'),
      Text(1, 0, '64'),
      Text(2, 0, '128'),
      Text(3, 0, '256'),
      Text(4, 0, '512')]
```

3.2 Max_features

The number of features to consider when looking for the best split.

```
[ ]: # Setting the max_features range
     max_features_list = [2, 5, 10, 50, 100, 200, 500, 1000, 2000, 5000, 10000,␣
      ↪20000]
     trs_list = list()
     cvs_list = list()

     for max_features in max_features_list:
         # Define model for each max_features
         rf_model = RandomForestClassifier(max_features=max_features)
         rf_model.fit(X_train, y_train)

         # Calculate the cross validation score
         train_score = accuracy_score(y_train, rf_model.predict(X_train))
         cv_score = np.mean(cross_val_score(rf_model, X_train, y_train, cv=5,␣
      ↪n_jobs=-1))

         trs_list.append(train_score)
         cvs_list.append(cv_score)
```

```
[ ]:  # Draw the plot for max_features
      fig = sns.lineplot(x=list(range(len(max_features_list))), y=cvs_list)
      fig = sns.lineplot(x=list(range(len(max_features_list))), y=trs_list)
      fig.set_xticks(range(len(max_features_list)))
      fig.set_xticklabels(max_features_list)
```

c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):

```
[ ]:  [Text(0, 0, '2'),
       Text(1, 0, '5'),
       Text(2, 0, '10'),
       Text(3, 0, '50'),
       Text(4, 0, '100'),
       Text(5, 0, '200'),
       Text(6, 0, '500'),
       Text(7, 0, '1000'),
       Text(8, 0, '2000'),
       Text(9, 0, '5000'),
       Text(10, 0, '10000'),
       Text(11, 0, '20000')]
```

3.3 Max_depth

max_depth is the maximum depth of the tree.

```
# Setting the possible value for max depth
max_depth_list = [20, 50, 100, 200, 500, 1000, 2000, 5000, 10000, 15000]

trs_list = list()
cvs_list = list()

for max_depth in max_depth_list:
    # Define model for each max_depth
    rf_model = RandomForestClassifier(max_depth=max_depth)
    rf_model.fit(X_train, y_train)

    # Calculate the cross validation score
    train_score = accuracy_score(y_train, rf_model.predict(X_train))
    cvs_score = np.mean(cross_val_score(rf_model, X_train, y_train, cv=5,
    n_jobs=-1))

    trs_list.append(train_score)
    cvs_list.append(cvs_score)
```

```python
# Draw the plot for max depth
fig = sns.lineplot(x=list(range(len(max_depth_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(max_depth_list))), y=trs_list)
fig.set_xticks(range(len(max_depth_list)))
fig.set_xticklabels(max_depth_list)
```

c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):

```
[ ]: [Text(0, 0, '20'),
      Text(1, 0, '50'),
      Text(2, 0, '100'),
      Text(3, 0, '200'),
      Text(4, 0, '500'),
      Text(5, 0, '1000'),
      Text(6, 0, '2000'),
      Text(7, 0, '5000'),
      Text(8, 0, '10000'),
      Text(9, 0, '15000')]
```

3.4 Min_samples_split

min_samples_split is the minimum number of samples required to split an internal node.

```python
# Setting the possible value for min_samples_split
min_samples_split_list = [10, 25, 50, 100, 200, 500, 1000, 2000, 5000]

trs_list = list()
cvs_list = list()

for min_samples_split in min_samples_split_list:
    # Define model for each min_samples_split
    rf_model = RandomForestClassifier(min_samples_split=min_samples_split)
    rf_model.fit(X_train, y_train)

    # Calculate the cross validation score
    train_score = accuracy_score(y_train, rf_model.predict(X_train))
    cvs_score = np.mean(cross_val_score(rf_model, X_train, y_train, cv=5,
    ↪n_jobs=-1))

    trs_list.append(train_score)
    cvs_list.append(cvs_score)
```

```python
# Draw the plot for min_samples_split
fig = sns.lineplot(x=list(range(len(min_samples_split_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(min_samples_split_list))), y=trs_list)
fig.set_xticks(range(len(min_samples_split_list)))
fig.set_xticklabels(min_samples_split_list)
```

c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):

```
[ ]: [Text(0, 0, '10'),
      Text(1, 0, '25'),
      Text(2, 0, '50'),
      Text(3, 0, '100'),
      Text(4, 0, '200'),
      Text(5, 0, '500'),
      Text(6, 0, '1000'),
      Text(7, 0, '2000'),
      Text(8, 0, '5000')]
```

3.5 Min_samples_leaf

min_samples_leaf is the minimum number of samples required to be at a leaf node. A split point at any depth will only be considered if it leaves at least ***min_samples_leaf*** training samples in each of the left and right branches.

```python
# Setting the min_samples_leaf range
min_samples_leaf_list = [1, 5, 10, 25, 50, 75, 100]
trs_list = list()
cvs_list = list()

for min_samples_leaf in min_samples_leaf_list:
    # Define model for each min_samples_leaf
    rf_model = RandomForestClassifier(min_samples_leaf=min_samples_leaf)
    rf_model.fit(X_train, y_train)

    # Calculate the cross validation score
    train_score = accuracy_score(y_train, rf_model.predict(X_train))
    cv_score = np.mean(cross_val_score(rf_model, X_train, y_train, cv=5,␣
    ↪n_jobs=-1))

    trs_list.append(train_score)
```

```
    cvs_list.append(cv_score)
```

```python
# Draw the plot for min_samples_split
fig = sns.lineplot(x=list(range(len(min_samples_leaf_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(min_samples_leaf_list))), y=trs_list)
fig.set_xticks(range(len(min_samples_leaf_list)))
fig.set_xticklabels(min_samples_leaf_list)
```

```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```

```
[ ]: [Text(0, 0, '1'),
      Text(1, 0, '5'),
      Text(2, 0, '10'),
      Text(3, 0, '25'),
      Text(4, 0, '50'),
      Text(5, 0, '75'),
      Text(6, 0, '100')]
```

4 Multiple tuning

First, we use grid search to help tuning this model.

```
[ ]: dict_param = {
         'max_depth' : np.asarray([500, 1000, 2000]),
         'min_samples_split': np.asarray([10, 200, 1000]),
         'min_samples_leaf': np.asarray([1, 5, 10]),
         'max_features': np.asarray([50, 200, 1000]),
     }

     grid_search = GridSearchCV(RandomForestClassifier(n_estimators=256),␣
      ↪dict_param, cv = 5, n_jobs=5)
     grid_search.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5, estimator=RandomForestClassifier(n_estimators=256), n_jobs=5,
                  param_grid={'max_depth': array([ 100,  500, 2000]),
                              'max_features': array([  50,  200, 1000]),
                              'min_samples_leaf': array([ 1,  5, 10]),
                              'min_samples_split': array([  10,  200, 2000])})
```

We elminate all parameters that appear in models with the validation accuracy < 0.85

```
df = pd.DataFrame(
  dict(
    max_depth = [val['max_depth'] for val in grid_search.cv_results_['params']],
    min_samples_split = [val['min_samples_split'] for val in grid_search.
 ↪cv_results_['params']],
    min_samples_leaf = [val['min_samples_leaf'] for val in grid_search.
 ↪cv_results_['params']],
    max_features = [val['max_features'] for val in grid_search.
 ↪cv_results_['params']],
    score = grid_search.cv_results_['mean_test_score']
  )
)

df = df[df['score'] <= 0.85]

for param in dict_param:
  for value in dict_param[param]:
    if len(df[df[param] == value]) == 81 // len(dict_param[param]) :
      print(param, value)
```

```
max_depth 100
min_samples_leaf 5
min_samples_leaf 10
max_features 1000
```

We repeat this process again, this time with the domain narrowed down.

```
dict_param = {
    'max_depth' : np.asarray([1000, 2000, 5000]),
    'min_samples_split': np.asarray([25, 200, 1000]),
    'min_samples_leaf': np.arange(1, 4),
    'max_features': np.asarray([50, 100, 200]),
}

grid_search = GridSearchCV(RandomForestClassifier(n_estimators=256),␣
 ↪dict_param, cv = 5, n_jobs=5)
grid_search.fit(X_train, y_train)
```

```
GridSearchCV(cv=5, estimator=RandomForestClassifier(n_estimators=256), n_jobs=5,
             param_grid={'max_depth': array([1000, 2000, 5000]),
                         'max_features': array([ 50, 100, 200]),
                         'min_samples_leaf': array([1, 2, 3]),
                         'min_samples_split': array([  25,  200, 1000])})
```

```
df = pd.DataFrame(
  dict(
    max_depth = [val['max_depth'] for val in grid_search.cv_results_['params']],
```

```
      min_samples_split = [val['min_samples_split'] for val in grid_search.
  ↪cv_results_['params']],
      min_samples_leaf = [val['min_samples_leaf'] for val in grid_search.
  ↪cv_results_['params']],
      max_features = [val['max_features'] for val in grid_search.
  ↪cv_results_['params']],
      score = grid_search.cv_results_['mean_test_score']
  )
)
df = df[df['score'] <= 0.85]

for param in dict_param:
  for value in dict_param[param]:
    if len(df[df[param] == value]) == 81 // len(dict_param[param]) :
      print(param, value)
```

Find the best combination of parameters for the model:

```
[ ]: print(grid_search.best_estimator_, grid_search.best_score_)
```

```
RandomForestClassifier(max_depth=1000, max_features=100, min_samples_split=25,
                       n_estimators=256) 0.8688749999999998
```

5   Conclusion

We use all the parameters from the last section to define the best model and then evaluate it using the preset functions.

```
[ ]: best_rf_model = RandomForestClassifier(max_depth=1000, max_features=100,␣
  ↪min_samples_split=25,
                         n_estimators=256)
  best_rf_model.fit(X_train, y_train)

  evaluate_model(best_rf_model, X_train, X_test, y_train, y_test,␣
  ↪include_training=True)
```

```
Score of on train are:
        - Accuracy score: 0.9651
        - Micro F1 score: 0.9651
        - Macro F1 score: 0.9569
Score of on test are:
        - Accuracy score: 0.8722
        - Micro F1 score: 0.8722
        - Macro F1 score: 0.8295
```

Confusion matrix on Train

Confusion matrix on Test

After that, we draw the learning curve of this Random forest model.

```
draw_learning_curve(best_rf_model, X_train, y_train)
```

Finally, we export the model.

```
directory = "data/models/rf/"

dump(best_rf_model, directory + "best_rf_model_bow_l1.joblib")
```

```
['data/models/rf/best_rf_model_bow_l1.joblib']
```