

Softmax Regression - tfidf

May 12, 2024

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import warnings
warnings.filterwarnings('ignore')

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from joblib import dump, load

from preset_function import evaluate_model, draw_learning_curve, \
    load_processed_data

X_train_bow, X_test_bow, X_train_tfidf, X_test_tfidf, \
    X_train_bow_L1, X_test_bow_L1, X_train_tfidf_L1, X_test_tfidf_L1 = \
    load_processed_data('input')

y_train, y_test = load_processed_data('output')

%matplotlib inline
```

Select dataset:

```
[ ]: X_train = X_train_tfidf
X_test = X_test_tfidf
```

1 Basic training

```
[ ]: softmax_model = LogisticRegression(multi_class='multinomial')
softmax_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(multi_class='multinomial')
```


```
[ ]: evaluate_model(softmax_model, X_train, X_test, y_train, y_test,
    ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9442
- Micro F1 score: 0.9442
- Macro F1 score: 0.9143


Score of on test are:

- Accuracy score: 0.8253
- Micro F1 score: 0.8253
- Macro F1 score: 0.7504





```
[ ]: draw_learning_curve(softmax_model, X_train, y_train)
```



2 Multiple tuning

2.1 No regularization

```
[ ]: softmax_model = LogisticRegression(penalty=None, solver='lbfgs',
    ↪ multi_class='multinomial')
softmax_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(multi_class='multinomial', penalty=None)
```


```
[ ]: evaluate_model(softmax_model, X_train, X_test, y_train, y_test,
    ↪ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9924
- Micro F1 score: 0.9924
- Macro F1 score: 0.9903


Score of on test are:

- Accuracy score: 0.8675
- Micro F1 score: 0.8675
- Macro F1 score: 0.8305





```
[ ]: draw_learning_curve(softmax_model, X_train, y_train)
```



2.2 L1 regularization

First, we try to plot the validation score through a list of C from 0.001 to 100

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l1', solver='saga',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
```

```
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)


# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[0.001, 0.01, 0.1, 1, 5, 10, 100]
```

```
[0.3386875, 0.3386875, 0.5438125, 0.88425, 0.971375, 0.992625, 0.992625]
```

```
[0.329, 0.33868750000000001, 0.40318749999999999, 0.86343750000000001,
0.87725000000000001, 0.875875, 0.870125]
```

```
[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



We can see the good value of C is near $C = 5$, then we scope to $C = 5$:

```
[ ]: C_list = [2, 3, 4, 4.5, 5, 5.5, 6, 7]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l1', solver='saga',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
    cvs_list.append(cv_score)
```


```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[2, 3, 4, 4.5, 5, 5.5, 6, 7]
[0.904625, 0.9281875, 0.9533125, 0.9629375, 0.97125, 0.9780625, 0.9829375,
0.989375]
[0.8749375, 0.8763124999999998, 0.8772499999999999, 0.8770625000000001,
0.8773125, 0.8768750000000001, 0.8763750000000001, 0.8758125]
```

```
[ ]: [Text(0, 0, '2'),
      Text(1, 0, '3'),
      Text(2, 0, '4'),
      Text(3, 0, '4.5'),
      Text(4, 0, '5'),
      Text(5, 0, '5.5'),
```

```
Text(6, 0, '6'),  
Text(7, 0, '7')]
```



We choose $C = 5$ to be the best model.

```
[ ]: best_l1_softmax_model = LogisticRegression(C=5, penalty='l1', solver='saga',  
↳ multi_class='multinomial')
```


```
[ ]: best_l1_softmax_model.fit(X_train, y_train)  
evaluate_model(best_l1_softmax_model, X_train, X_test, y_train, y_test,  
↳ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9714
- Micro F1 score: 0.9714
- Macro F1 score: 0.9621


Score of on test are:

- Accuracy score: 0.8808
- Micro F1 score: 0.8808
- Macro F1 score: 0.8493





```
[ ]: draw_learning_curve(best_l1_softmax_model, X_train, y_train)
```



2.3 L2 regularization

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100, 200]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))


    trs_list.append(train_score)
    cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)

[0.001, 0.01, 0.1, 1, 5, 10, 100, 200]
[0.3386875, 0.381625, 0.6619375, 0.9441875, 0.99125, 0.9920625, 0.992625,
0.9921875]
[0.33868750000000001, 0.36025, 0.6044375, 0.811625, 0.84781250000000001, 0.8535,
0.86075, 0.8615]

[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100'),
      Text(7, 0, '200')]
```



We can see the good value of C is near $C = 100$, then we scope to $C = 100$:

```
[ ]: C_list = [90, 95, 98, 99, 100, 101, 102, 105, 110]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
    cvs_list.append(cv_score)
```


```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[90, 95, 98, 99, 100, 101, 102, 105, 110]
[0.992625, 0.992625, 0.992125, 0.99225, 0.992625, 0.992625, 0.992625, 0.992625,
0.9925625]
[0.8595624999999998, 0.8609375, 0.8618124999999999, 0.8612499999999998, 0.86075,
0.8590625, 0.8603124999999998, 0.8618750000000001, 0.8619375]
```

```
[ ]: [Text(0, 0, '90'),
      Text(1, 0, '95'),
      Text(2, 0, '98'),
      Text(3, 0, '99'),
      Text(4, 0, '100'),
      Text(5, 0, '101'),
```

```
Text(6, 0, '102'),
Text(7, 0, '105'),
Text(8, 0, '110')]
```



We choose $C = 110$ to be the best model.

```
[ ]: best_l2_softmax_model = LogisticRegression(C=110, penalty='l2', solver='lbfgs',
↳multi_class='multinomial')
```


```
[ ]: best_l2_softmax_model.fit(X_train, y_train)
evaluate_model(best_l2_softmax_model, X_train, X_test, y_train, y_test,
↳include_training=True)
```


Score of on train are:

- Accuracy score: 0.9926
- Micro F1 score: 0.9926
- Macro F1 score: 0.9904


Score of on test are:

- Accuracy score: 0.8692
- Micro F1 score: 0.8692
- Macro F1 score: 0.8304





```
[ ]: draw_learning_curve(best_l2_softmax_model, X_train, y_train)
```



2.4 Elastic regularization

```
[ ]: dict_param = {
    'C' : [0.001, 0.01, 0.1, 1, 5, 10, 100],
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

softmax_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='multinomial')
grid_search = GridSearchCV(softmax_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='multinomial',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': [0.001, 0.01, 0.1, 1, 5, 10, 100],
                'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```

```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
df = df[df['score'] < 0.8]
print("Bad hyperparameter:")
for param in dict_param:
    for value in dict_param[param]:
        if len(df[df[param] == value]) == 35 // len(dict_param[param]):
            print(param, value)
```

Bad hyperparameter:

C 0.001

C 0.01

C 0.1

```
[ ]: dict_param = {
    'C' : np.logspace(0, 2, 5),
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

softmax_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='multinomial')
grid_search = GridSearchCV(softmax_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='multinomial',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': array([ 1.          ,  3.16227766, 10.          ,
31.6227766 ,
100.          ]),
    'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```

```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
print(df)
```

	C	l1_ratio	score
0	1.000000	0.1	0.821063
1	1.000000	0.3	0.831812
2	1.000000	0.5	0.839813
3	1.000000	0.7	0.848938
4	1.000000	0.9	0.861063
5	3.162278	0.1	0.849187
6	3.162278	0.3	0.859062
7	3.162278	0.5	0.866500
8	3.162278	0.7	0.871562
9	3.162278	0.9	0.876250
10	10.000000	0.1	0.857062
11	10.000000	0.3	0.864437
12	10.000000	0.5	0.869438
13	10.000000	0.7	0.872188
14	10.000000	0.9	0.875313
15	31.622777	0.1	0.858938
16	31.622777	0.3	0.864375
17	31.622777	0.5	0.867563
18	31.622777	0.7	0.869813
19	31.622777	0.9	0.872750
20	100.000000	0.1	0.860750
21	100.000000	0.3	0.862750
22	100.000000	0.5	0.865313
23	100.000000	0.7	0.867750
24	100.000000	0.9	0.868375

```
[ ]: print(grid_search.best_estimator_, grid_search.best_score_)
```

```
LogisticRegression(C=3.1622776601683795, l1_ratio=0.9,
                    multi_class='multinomial', penalty='elasticnet',
                    solver='saga') 0.87625
```

```
[ ]: best_en_softmax_model = LogisticRegression(C=3.1622776601683795, l1_ratio=0.9,
                    multi_class='multinomial', penalty='elasticnet',
                    solver='saga')
```


```
[ ]: best_en_softmax_model.fit(X_train, y_train)
evaluate_model(best_en_softmax_model, X_train, X_test, y_train, y_test,
               include_training=True)
```


Score of on train are:

- Accuracy score: 0.9329
- Micro F1 score: 0.9329
- Macro F1 score: 0.9098


Score of on test are:

- Accuracy score: 0.8810
- Micro F1 score: 0.8810
- Macro F1 score: 0.8473





```
[ ]: draw_learning_curve(best_en_softmax_model, X_train, y_train)
```



3 Conclusion

There are a few difference among the accuracy of these 3 regularization. However, Elastic-net regularization gives the best performance then I will choose it to be the best model in this notebook.

```
[ ]: best_softmax_model = best_en_softmax_model
```

```
[ ]: directory = "data/models/softmax/"  
  
     dump(best_softmax_model, directory + "best_softmax_tfidf_model.joblib")
```

```
[ ]: ['data/models/softmax/best_softmax_tfidf_model.joblib']
```