

# Softmax Regression - BoW

May 12, 2024

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import warnings
warnings.filterwarnings('ignore')

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from joblib import dump, load

from preset_function import evaluate_model, draw_learning_curve, \
    load_processed_data

X_train_bow, X_test_bow, X_train_tfidf, X_test_tfidf, \
    X_train_bow_L1, X_test_bow_L1, X_train_tfidf_L1, X_test_tfidf_L1 = \
    load_processed_data('input')

y_train, y_test = load_processed_data('output')

%matplotlib inline
```

Select dataset:

```
[ ]: X_train = X_train_bow
X_test = X_test_bow
```

## 1 Basic training

```
[ ]: softmax_model = LogisticRegression(multi_class='multinomial')
softmax_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(multi_class='multinomial')
```


```
[ ]: evaluate_model(softmax_model, X_train, X_test, y_train, y_test,
    ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9889
- Micro F1 score: 0.9889
- Macro F1 score: 0.9866


Score of on test are:

- Accuracy score: 0.8740
- Micro F1 score: 0.8740
- Macro F1 score: 0.8371





```
[ ]: draw_learning_curve(softmax_model, X_train, y_train)
```



## 2 Multiple tuning

### 2.1 No regularization

```
[ ]: softmax_model = LogisticRegression(penalty=None, solver='lbfgs',
    ↪ multi_class='multinomial')
softmax_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(multi_class='multinomial', penalty=None)
```


```
[ ]: evaluate_model(softmax_model, X_train, X_test, y_train, y_test,
    ↪ include_training=True)
```


Score of on train are:

- Accuracy score: 0.9925
- Micro F1 score: 0.9925
- Macro F1 score: 0.9905


Score of on test are:

- Accuracy score: 0.8620
- Micro F1 score: 0.8620
- Macro F1 score: 0.8282





```
[ ]: draw_learning_curve(softmax_model, X_train, y_train)
```



## 2.2 L1 regularization

First, we try to plot the validation score through a list of C from 0.001 to 100

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l1', solver='saga',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
```

```
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)


# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[0.001, 0.01, 0.1, 1, 5, 10, 100]
```

```
[0.3386875, 0.338625, 0.859, 0.926125, 0.987625, 0.9890625, 0.99025]
```

```
[0.33868750000000001, 0.33831249999999996, 0.836, 0.877, 0.8730625,
0.87237500000000001, 0.8703125]
```

```
[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```





We can see the good value of  $C$  is near  $C = 1$ , then we scope to  $C = 1$ :

```
[ ]: C_list = [0.1, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l1', solver='saga',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
    cvs_list.append(cv_score)
```


```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[0.1, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2]
[0.8589375, 0.88675, 0.903625, 0.9145625, 0.926125, 0.9350625, 0.94225,
0.9521875, 0.9649375]
[0.83575, 0.8668125, 0.8762500000000001, 0.8771875, 0.8768750000000001,
0.8759375, 0.8748750000000001, 0.874375, 0.8730625]
```

```
[ ]: [Text(0, 0, '0.1'),
      Text(1, 0, '0.25'),
      Text(2, 0, '0.5'),
      Text(3, 0, '0.75'),
      Text(4, 0, '1'),
      Text(5, 0, '1.25'),
```

```
Text(6, 0, '1.5'),
Text(7, 0, '1.75'),
Text(8, 0, '2')]
```



We choose  $C = 1$  to be the best model.

```
[ ]: best_l1_softmax_model = LogisticRegression(C=1, penalty='l1', solver='saga',
↳multi_class='multinomial')
```


```
[ ]: best_l1_softmax_model.fit(X_train, y_train)
evaluate_model(best_l1_softmax_model, X_train, X_test, y_train, y_test,
↳include_training=True)
```


Score of on train are:

- Accuracy score: 0.9263
- Micro F1 score: 0.9263
- Macro F1 score: 0.9073


Score of on test are:

- Accuracy score: 0.8802
- Micro F1 score: 0.8802
- Macro F1 score: 0.8501





```
[ ]: draw_learning_curve(best_l1_softmax_model, X_train, y_train)
```



## 2.3 L2 regularization

We do the same with L1 regularization

```
[ ]: C_list = [0.001, 0.01, 0.1, 1, 5, 10, 100]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
```

```
cvs_list.append(cv_score)
```

```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)


# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[0.001, 0.01, 0.1, 1, 5, 10, 100]
```

```
[0.390375, 0.6790625, 0.94425, 0.9889375, 0.992375, 0.992625, 0.9925625]
```

```
[0.360625, 0.60500000000000001, 0.84406249999999999, 0.86856250000000001, 0.8665,
0.8655625, 0.860375]
```

```
[ ]: [Text(0, 0, '0.001'),
      Text(1, 0, '0.01'),
      Text(2, 0, '0.1'),
      Text(3, 0, '1'),
      Text(4, 0, '5'),
      Text(5, 0, '10'),
      Text(6, 0, '100')]
```



We can see the good value of  $C$  is near  $C = 1$ , then we scope to  $C = 1$ :

```
[ ]: C_list = [0.1, 0.25, 0.5, 1, 1.25, 1.5, 1.75, 2]

# Define a list in order to store accuracy points
cvs_list = list()
trs_list = list()

for c in C_list:
    # Define model for each C
    softmax_model = LogisticRegression(C=c, penalty='l2', solver='lbfgs',
    ↪ multi_class='multinomial')
    softmax_model.fit(X_train, y_train)

    # Calculate score of cross validation
    train_score = accuracy_score(y_train, softmax_model.predict(X_train))
    cv_score = np.mean(cross_val_score(softmax_model, X_train, y_train, cv=5,
    ↪ n_jobs=8))

    trs_list.append(train_score)
    cvs_list.append(cv_score)
```


```
[ ]: # Print the result
print(C_list)
print(trs_list)
print(cvs_list)

# Draw the plot
fig = sns.lineplot(x=list(range(len(C_list))), y=cvs_list)
fig = sns.lineplot(x=list(range(len(C_list))), y=trs_list)
fig.set_xticks(range(len(C_list)))
fig.set_xticklabels(C_list)
```

```
[0.1, 0.25, 0.5, 1, 1.25, 1.5, 1.75, 2]
[0.94425, 0.971375, 0.9826875, 0.9889375, 0.9898125, 0.99075, 0.991125,
0.9914375]
[0.8440624999999999, 0.8591875, 0.8660625, 0.8685625000000001, 0.8684375,
0.8681875, 0.867875, 0.8678125]
```

```
[ ]: [Text(0, 0, '0.1'),
      Text(1, 0, '0.25'),
      Text(2, 0, '0.5'),
      Text(3, 0, '1'),
      Text(4, 0, '1.25'),
      Text(5, 0, '1.5'),
```

```
Text(6, 0, '1.75'),  
Text(7, 0, '2')]
```



We choose  $C = 1$  to be the best model.

```
[ ]: best_l2_softmax_model = LogisticRegression(C=1, penalty='l2', solver='lbfgs',  
↪multi_class='multinomial')
```

```
[ ]: best_l2_softmax_model.fit(X_train, y_train)  
evaluate_model(best_l2_softmax_model, X_train, X_test, y_train, y_test,  
↪include_training=True)
```


Score of on train are:


- Accuracy score: 0.9889
- Micro F1 score: 0.9889
- Macro F1 score: 0.9866

Score of on test are:


- Accuracy score: 0.8740
- Micro F1 score: 0.8740
- Macro F1 score: 0.8371







```
[ ]: draw_learning_curve(best_l2_softmax_model, X_train, y_train)
```



## 2.4 Elastic regularization

```
[ ]: dict_param = {
    'C' : [0.001, 0.01, 0.1, 1, 5, 10, 100],
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

softmax_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='multinomial')
grid_search = GridSearchCV(softmax_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='multinomial',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': [0.001, 0.01, 0.1, 1, 5, 10, 100],
                'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```

```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
df = df[df['score'] < 0.8]
print("Bad hyperparameter:")
for param in dict_param:
    for value in dict_param[param]:
        if len(df[df[param] == value]) == 35 // len(dict_param[param]):
            print(param, value)
```

Bad hyperparameter:

C 0.001

C 0.01

```
[ ]: dict_param = {
    'C' : np.logspace(0, 2, 5),
    'l1_ratio' : np.linspace(0.1, 0.9, 5)
}

softmax_model = LogisticRegression(penalty='elasticnet', solver='saga',
    ↪multi_class='multinomial')
grid_search = GridSearchCV(softmax_model, dict_param, scoring='accuracy', cv=5,
    ↪n_jobs=-1)
grid_search.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5,
    estimator=LogisticRegression(multi_class='multinomial',
                                penalty='elasticnet', solver='saga'),
    n_jobs=-1,
    param_grid={'C': array([ 1.          ,  3.16227766, 10.          ,
31.6227766 ,
100.          ]),
    'l1_ratio': array([0.1, 0.3, 0.5, 0.7, 0.9])},
    scoring='accuracy')
```

```
[ ]: df = pd.DataFrame(
    dict(
        C = [val['C'] for val in grid_search.cv_results_['params']],
        l1_ratio = [val['l1_ratio'] for val in grid_search.cv_results_['params']],
        score = grid_search.cv_results_['mean_test_score']
    )
)
print(df)
```

	C	l1_ratio	score
0	1.000000	0.1	0.871438
1	1.000000	0.3	0.873062
2	1.000000	0.5	0.874437
3	1.000000	0.7	0.875188
4	1.000000	0.9	0.876250
5	3.162278	0.1	0.870250
6	3.162278	0.3	0.872000
7	3.162278	0.5	0.873125
8	3.162278	0.7	0.873000
9	3.162278	0.9	0.872812
10	10.000000	0.1	0.869750
11	10.000000	0.3	0.870562
12	10.000000	0.5	0.871125
13	10.000000	0.7	0.871250
14	10.000000	0.9	0.871562
15	31.622777	0.1	0.869875
16	31.622777	0.3	0.870188
17	31.622777	0.5	0.870187
18	31.622777	0.7	0.870125
19	31.622777	0.9	0.870125
20	100.000000	0.1	0.869250
21	100.000000	0.3	0.869938
22	100.000000	0.5	0.870188
23	100.000000	0.7	0.870125
24	100.000000	0.9	0.870187

```
[ ]: print(grid_search.best_estimator_, grid_search.best_score_)
```

```
LogisticRegression(l1_ratio=0.9, multi_class='multinomial',
                    penalty='elasticnet', solver='saga') 0.87625
```

```
[ ]: best_en_softmax_model = LogisticRegression(C=1, l1_ratio=0.9,
        ↪multi_class='multinomial',
        ↪penalty='elasticnet', solver='saga')
```


```
[ ]: best_en_softmax_model.fit(X_train, y_train)
evaluate_model(best_en_softmax_model, X_train, X_test, y_train, y_test,
        ↪include_training=True)
```


Score of on train are:

- Accuracy score: 0.9293
- Micro F1 score: 0.9293
- Macro F1 score: 0.9113


Score of on test are:

- Accuracy score: 0.8808
- Micro F1 score: 0.8808
- Macro F1 score: 0.8505





```
[ ]: draw_learning_curve(best_en_softmax_model, X_train, y_train)
```



### 3 Conclusion

There are a few difference among the accuracy of these 3 regularization. However, Elastic-net regularization gives the best performance then I will choose it to be the best model in this notebook.

```
[ ]: best_softmax_model = best_en_softmax_model
```

```
[ ]: directory = "data/models/softmax/"  
     dump(best_softmax_model, directory + "best_softmax_bow_model.joblib")
```

```
[ ]: ['data/models/softmax/best_softmax_bow_model.joblib']
```